

# Final Project Guidelines

## Project Overview

This capstone project requires you to design and build a **functional AI application** as a team. You will create a production-ready web application that addresses a real-world problem or introduces an innovative solution using LLMs and the technical skills developed throughout the course.

Each team will develop a unique application based on a problem or opportunity you have identified. This project demonstrates your ability to:

- Design and implement AI-powered applications
- Work collaboratively in a team environment
- Apply professional software architecture patterns
- Deploy and monitor production AI systems
- Transform concepts into working software

**Important:** See the Deliverables section for submission deadlines and requirements.

## Technical Requirements

### Core Requirements

#### Backend: Python (Required)

- Python is the industry standard for AI application development
- Must integrate with an LLM API
- Recommended: Google Gemini (free tier available)
- Alternatives: OpenAI GPT, Anthropic Claude, or other major providers (with course professor approval)

#### Frontend: Your Choice

- Streamlit (covered in class) - simplest option for rapid development
- React, Vue, Next.js, or other modern frameworks
- Any technology that suits your project needs (with course professor approval)

#### AI Observability & Monitoring

- Langfuse (covered in class) or alternative observability tools (with course professor approval)
- Required for debugging and monitoring your AI application

#### Deployment

- Application must be publicly accessible
- Platform of your choice (Streamlit Cloud, Render, Vercel, AWS, etc.)

#### Documentation

- Root README.md with comprehensive project overview (see README\_Template.md for example structure)
- docs/ARCHITECTURE.md explaining architecture decisions and technical choices

- Component-level documentation recommended for complex modules
- Complete setup and deployment instructions

## Bonus Opportunities

Going beyond the core requirements is encouraged and rewarded. Examples include:

- **Building a REST API:** Professional API architecture demonstrating advanced development skills
- **Exceptional UI/UX:** Polished, intuitive interfaces that enhance user experience
- **Cross-disciplinary integration:** Incorporating machine learning models, predictive analytics, or concepts from other courses
- **Real-world data:** Using actual datasets rather than mock data (e.g., Portuguese market data for predictive models)
- **Other advanced features:** Implementing additional functionality that meaningfully enhances your application's value

**Key principle:** Innovation and extra effort are valued when they serve your project's goals. Always justify your technical decisions.

**Important:** Any technology choices outside the recommended stack require course professor approval before implementation. Reach out via email mcardoso@novaims.unl.pt or in person to discuss.

## Project Scope

### Minimum Viable Product (MVP)

Your application must demonstrate the skills taught in class. Design a project that naturally requires these techniques.

#### Required Components:

1. **AI Integration:** LLM integration with proper prompting and structured outputs when needed (e.g., extracting data, routing logic, function parameters)
2. **Conversational Interface:** Multi-turn dialogue with the AI
  - Users must be able to interact through chat-based conversation
  - Enables follow-up questions, exploring results, and refining outputs
  - Note: Very few projects genuinely don't benefit from conversational interaction. If you believe yours is an exception, reach out to discuss - you'll need strong justification
3. **Function Calling / Tools:** AI agent with custom tools that extend LLM capabilities
4. **Document Ingestion:** Process and utilize document content
  - Core feature (analyzing user documents) or supporting feature (company FAQs, knowledge base)
  - Must serve a purpose, not just exist to tick a box
5. **Clean Architecture:** Layered structure with separation of concerns
6. **Observability:** Tracing integrated throughout (Langfuse or approved alternative)

#### Recommended:

- **Multimodal processing:** If your project benefits from image/PDF analysis

#### Advanced (When It Adds Real Value):

- **RAG with vector databases:** For similarity matching and semantic search use cases (finding similar products, matching user queries to past support tickets, research paper recommendations). Note: With modern 200K+ token context windows, most document Q&A doesn't need RAG - you can process documents directly within context. Use RAG/vector search when you genuinely need "find similar items" functionality or have knowledge bases larger than context limits, not just to check a box.

**Key Principle:** Balance is critical. Don't add features just to meet requirements, but don't avoid techniques you've learned either. Equally important: don't overcomplicate features that work fine as-is. Apply techniques where they add value, keep things simple where they don't.

### Evaluation Criteria

#### What We're Looking For:

##### Technical Implementation

- Clean architecture with proper separation of concerns
- Effective use of AI capabilities (LLM integration, prompting, structured outputs)
- Function calling/tools that genuinely extend your application
- Document ingestion that serves a purpose
- Proper observability and tracing throughout

##### Functionality & User Experience

- Application works reliably and solves the stated problem
- Conversational interface enables meaningful interaction
- Intuitive, polished user experience
- Features make sense for the problem you're solving

##### Innovation & Effort

- Going beyond minimum requirements where it adds value
- Creative solutions to real problems
- Technical depth and sophistication
- Evidence of research and learning beyond class material

##### Documentation & Professionalism

- Industry-grade README with clear setup instructions
- Architecture explanation and justification of technical choices
- Clean, well-organized code
- Professional presentation and defense

##### Presentation & Defense

- Compelling investor pitch
- Effective demonstration of key features

- Ability to defend technical decisions
- Clear communication of project value

**Key Principle:**

Quality over quantity. A well-executed application that thoughtfully applies appropriate techniques will score higher than a feature-bloated project that awkwardly forces every technique without purpose.

**Deliverables****1. Technical Report**

**Due: December 7, 23:59**

Submit via Moodle a detailed technical report outlining your project plan and implementation details.

**Filename:** TeamName\_TechnicalReport.pdf

See the Technical Report Guidelines document for complete requirements.

**2. Final Project Submission**

**Due: December 22, 23:59**

**Moodle Submission:**

- Submit complete project as a zip file via Moodle
- **Filename:** TeamName\_FinalProject.zip

**GitHub Repository:**

- Add mb-cardoso to your project repository as a collaborator
- Repository must include:
  - Complete source code
  - Root README.md with comprehensive overview and setup instructions
  - docs/ARCHITECTURE.md with architecture decisions and technical justifications
  - Component-level documentation (recommended for complex modules)
  - Environment setup instructions (dependencies, API keys, etc.)
  - Clear documentation of how to run the application locally

**Deployed Application:**

- Application must be deployed and publicly accessible
- Include deployment URL in README
- Platform choice is yours (Streamlit Cloud, Render, Vercel, AWS, etc.)

**3. Project Defense**

**Date: January 17 (Saturday), in person, time slots will be attributed**

**Duration: 15-minute presentation + up to 15-minute Q&A (maximum 30 minutes total)**

- **Investor Pitch:** Present your application as a startup seeking investment (presentation format of your choice)

- **Live Demo:** Demonstrate core functionality and key features
- **Technical Q&A:** Defend technical decisions, architecture choices, and implementation details
- **Evaluation Panel:** Course professor + AI industry professional
- **Team Participation:** All team members must attend and be prepared to answer questions about any aspect of the project

## Getting Started

This is an overview of the entire project timeline, from initial conception to final defense.

### 1. Form Your Team & Choose Your Problem

- Form team composition
- Identify a real-world problem or innovation opportunity
- Brainstorm how AI can address this problem

### 2. Define Your Application Concept

- What problem does your application solve?
- Who are your users?
- What's the core value proposition?
- How will users interact with it?

### 3. Plan Your Architecture

- Review clean architecture materials from class
- Design your layered structure (UI, Service, AI, Tools)
- Identify which AI capabilities you need:
  - What tools/functions will extend the LLM?
  - What documents need to be processed?
  - Does your project benefit from multimodal processing?
  - Do you need RAG/vector database (e.g., semantic search for similarity matching)?

### 4. Start Building

- Set up project structure
- Build core functionality iteratively
- Integrate observability from the beginning
- Deploy early for testing and iteration

### 5. Technical Report Submission

- Document your architecture and implementation progress
- Justify your technical choices
- Outline remaining features and development plan

- See the Technical Report Guidelines document
- See Deliverables section for deadline and submission details

## 6. Continue Development Through December

- Complete remaining features
- Test and refine user experience
- Keep deployment updated
- Maintain clean, documented code

## 7. Final Submission

- Industry-grade README and documentation
- Complete, deployed application
- Clean codebase with proper setup instructions
- See Deliverables section for deadline and submission details

## 8. Project Defense

- Investor pitch presentation
- Live demonstration
- Technical Q&A defense
- See Deliverables section for date, format, and requirements

## Resources & Documentation

### Official Documentation

#### Google Gemini

- Google AI Studio
- Gemini API Documentation
- Google GenAI Python SDK

#### Streamlit

- Streamlit Documentation
- Streamlit Components
- Streamlit Cloud Deployment

#### Langfuse

- Langfuse Documentation
- Python SDK Integration
- Tracing Guide

## Alternative LLM Providers

### OpenAI

- OpenAI API Documentation

- Python SDK

#### **Anthropic Claude**

- Claude API Documentation
- Python SDK

#### **Additional Tools**

##### **UV Package Manager**

- UV Documentation

##### **Database Options** (if needed for your project)

- MongoDB Atlas + PyMongo
- PostgreSQL + psycopg2
- SQLite (built into Python, file-based)
- Local file-based storage (JSON, CSV, etc.)

##### **ChromaDB** (if using RAG/vector databases)

- ChromaDB Documentation

#### **Course Materials**

All course materials and examples are available on Moodle.

#### **Documentation Templates**

- README\_Template.md - Example README structure for your project